


MEMORY MANAGEMENT: SEGMENTATION AND PAGING

Kai Mast

CS 537

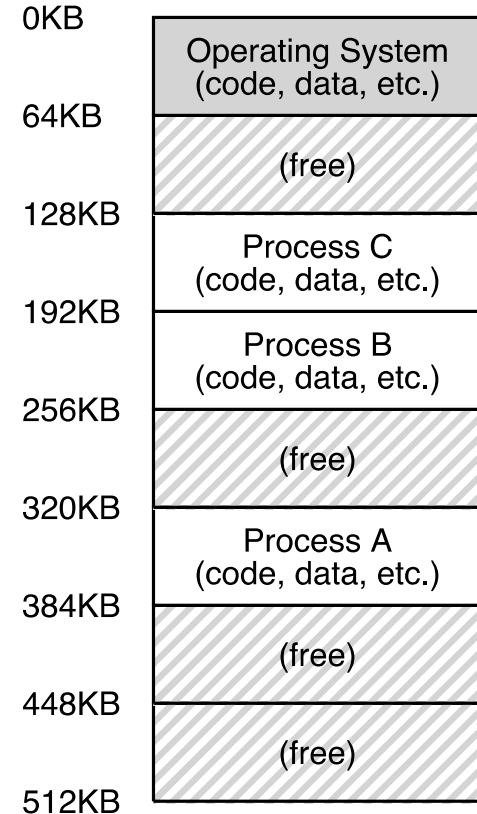
Fall 2022

RECAP: VIRTUAL MEMORY




Each process has own set of addresses

How can the OS provide illusion of private (virtual) address space to each process?



RECAP: TIME-SHARE MEMORY



- create P1
- swap out P1
- create P2
- swap out P2
- swap in P1
- ...

RECAP: STATIC RELOCATION


```
0x10:movl 0x8(%rbp), %edi  
0x13:addl $0x3, %edi  
0x19:movl %edi, 0x8(%rbp)
```

**Rewrite
for P1**



```
0x1010: movl 0x8(%rbp), %edi  
0x1013: addl $0x3, %edi  
0x1019: movl %edi, 0x8(%rbp)
```

**Rewrite
for P2**



```
0x3010: movl 0x8(%rbp), %edi  
0x3013: addl $0x3, %edi  
0x3019: movl %edi, 0x8(%rbp)
```

Process 1

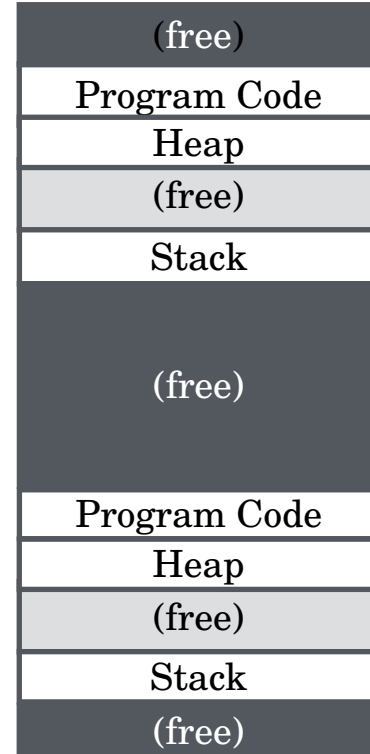
4 KB

8 KB

12 KB

Process 2

16 KB



DYNAMIC RELOCATION USING BASE REGISTERS

Base Registers

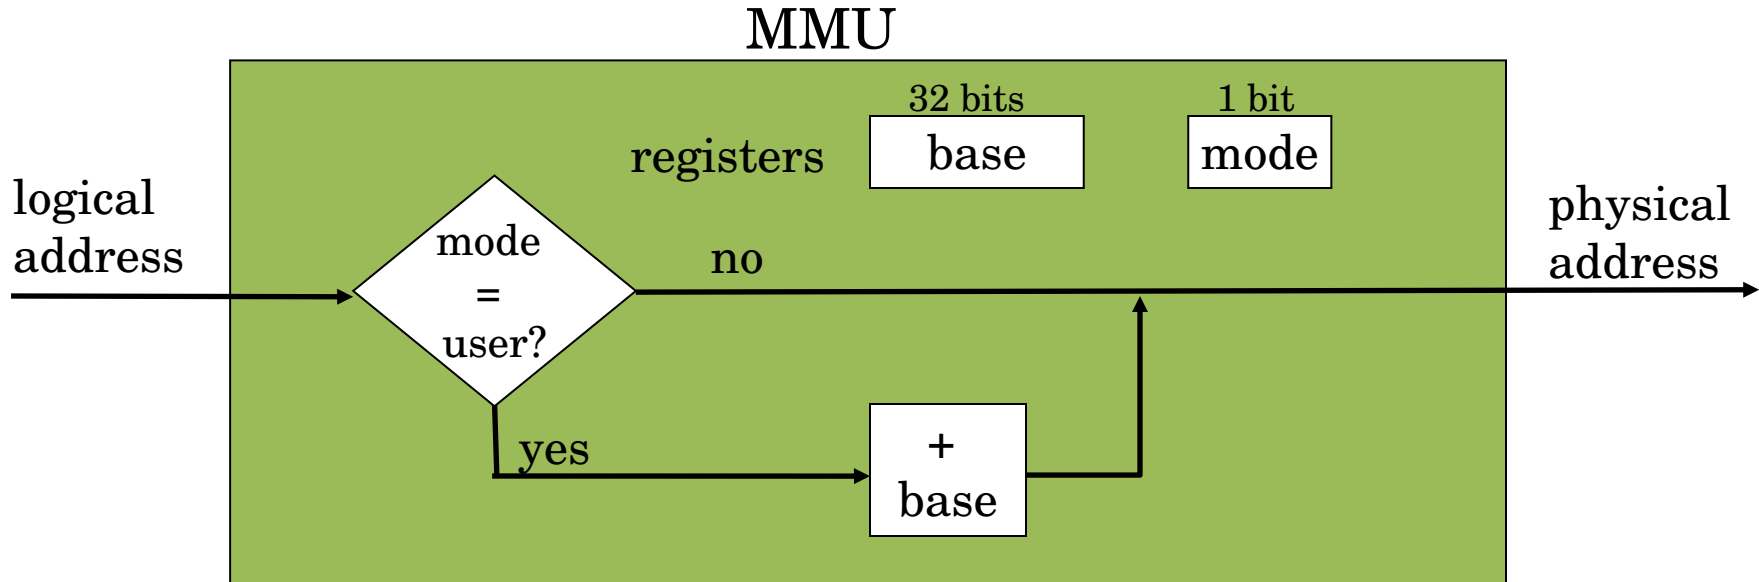
- Stored in the MMU
- Each process has different value in base register
- Set **by the OS** on **context switch**

Address Translation using Base Registers

- Add base register to virtual address on every memory access
- Dynamic relocation by changing value of base register!

BASE REGISTERS: IMPLEMENTATION

- Translation on every memory access of user process
- MMU adds base register to logical address to form physical address



4) DYNAMIC RELOCATION WITH BASE+BOUNDS

Idea: Limit the address space with a bounds register

Base register: Smallest physical address (or starting location)

Bounds register: Size of this process's virtual address space

- Sometimes defined as largest physical address (base + size)

If process loads/stores beyond bounds, what does the OS do? Kill process

What performs the bounds check? MMU

PROTECTION WITH BASE + BOUNDS



Possible Execution

Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5096, R1
P1: load 1000, R1	load 2024, R1
P1: store 3072, R1	kill P1

Process	Base	Bounds
P1	1024	1024
P2	4096	1024

MANAGING PROCESSES WITH BASE AND BOUNDS

Track base and bounds registers in PCB (process-control-block)

Context-Switch Steps

- OS in privileged mode
- Save base and bounds registers of old process
- Load base and bounds registers of new process
- Change to user mode and jump to new process

What if we do not change base and bounds registers during context switch?

Threads!

Protection Requirements

- User process cannot switch to privileged mode (except using traps)
- User process cannot change base and bounds registers

BASE AND BOUNDS SUMMARY

Advantages

- Provides protection (both read and write) across address spaces
- Supports dynamic relocation
- Can place process at different locations initially and move address spaces
- Simple, inexpensive implementation:
 - Few registers, little logic in MMU
- Fast: MMU can add and compare in parallel

Disadvantages

- Each process must be allocated contiguously in physical memory
- Must reserve memory that may not be used by process
- No partial sharing between processes
 - Cannot share limited parts of address space

5) SEGMENTATION

Divide logical address space into logical segments – visible to the OS


- Each segment corresponds to logical entity in address space (code, stack, heap)

Each segment has separate base + bounds register in the MMU

Advantages: Each segment can independently

- be placed separately in physical memory
- grow and shrink
- be protected (separate read/write/execute bits)

(Book Chapter 16)



SEGMENTED ADDRESSING

For virtual address, process specifies 1) segment and 2) offset within segment

How does process designate a particular segment?

- Use part of logical address
 - Top bits of logical address select segment
 - Low bits of logical address select offset within segment

SEGMENTATION IMPLEMENTATION

MMU contains Segment Table for the current process:

- Each segment has own base and bounds, protection bits
- Example: 14 bit logical address, 4 segments;

How many
bits for the
segment?

How many
bits for the
offset?

Segment	Base	Bounds	R W
0	0x2000	0x6ff	1 0
1	0x0000	0x4ff	1 1
2	0x3000	0xfff	1 1
3	0x0000	0x000	0 0

Remember:

A hex digit (4 bits)
has 16 possible values

ADDRESS TRANSLATIONS WITH SEGMENTATION

14 bit logical address, 4 segments;

Segment	Base	Bounds	R W
0	0x2000	0x6fff	1 0
1	0x0000	0x4fff	1 1
2	0x3000	0xffff	1 1
3	0x0000	0x0000	0 0

Remember:

1 hex digit 4 bits

Translate logical (in hex) to physical

0x0240:

0x1108:

0x265c:

0x3002:

ADVANTAGES OF SEGMENTATION

Enables sparse allocation of address space


Stack and heap can grow independently

- Heap: If no data on free list, dynamic memory allocator requests more from OS
(e.g., UNIX: malloc library makes sbrk() system call)
- Stack: OS recognizes reference near outside legal segment, extends stack implicitly

Different protection for different segments

- Enables sharing of selected segments
- Read-only status for code

Supports dynamic relocation of each segment




DISADVANTAGES OF SEGMENTATION

Each segment must be allocated contiguously

May not have sufficient physical memory for large segments?


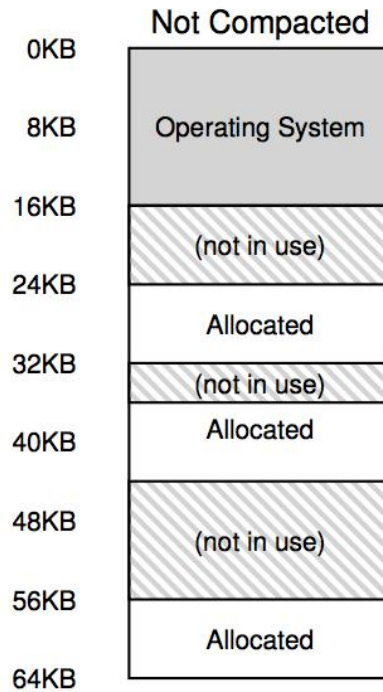
External Fragmentation
(memory fragmentation caused by the OS)



EXTERNAL FRAGMENTATION


Why is this bad?

Small pieces of free memory
might not be usable/assignable



INTERNAL FRAGMENTATION

What happens when you allocate memory?



Why is this additional layer needed?

- System calls are expensive
- VM is allocated at a coarse granularity (when using pages)

Why can the heap get fragmented?

- Allocation patterns are hard to predict
- malloc() may be called with vastly different sizes

REVIEW: MATCH DESCRIPTION

Description	Name of Approach
One process uses RAM at a time	Time Sharing
Rewrite code and addresses before running	Static Relocation
Add per-process starting location to virtual address to obtain phys address	Dynamic w/ Base
Dynamic approach that verifies address is in valid range	Dynamic w/ Base+Bounds
Several base+bound pairs per process	Segmentation

Candidates: Segmentation, Static Relocation, Base, Base+Bounds, Time Sharing

MEMORY MANAGEMENT WITH PAGING

(Book Chapter 18)

PAGING

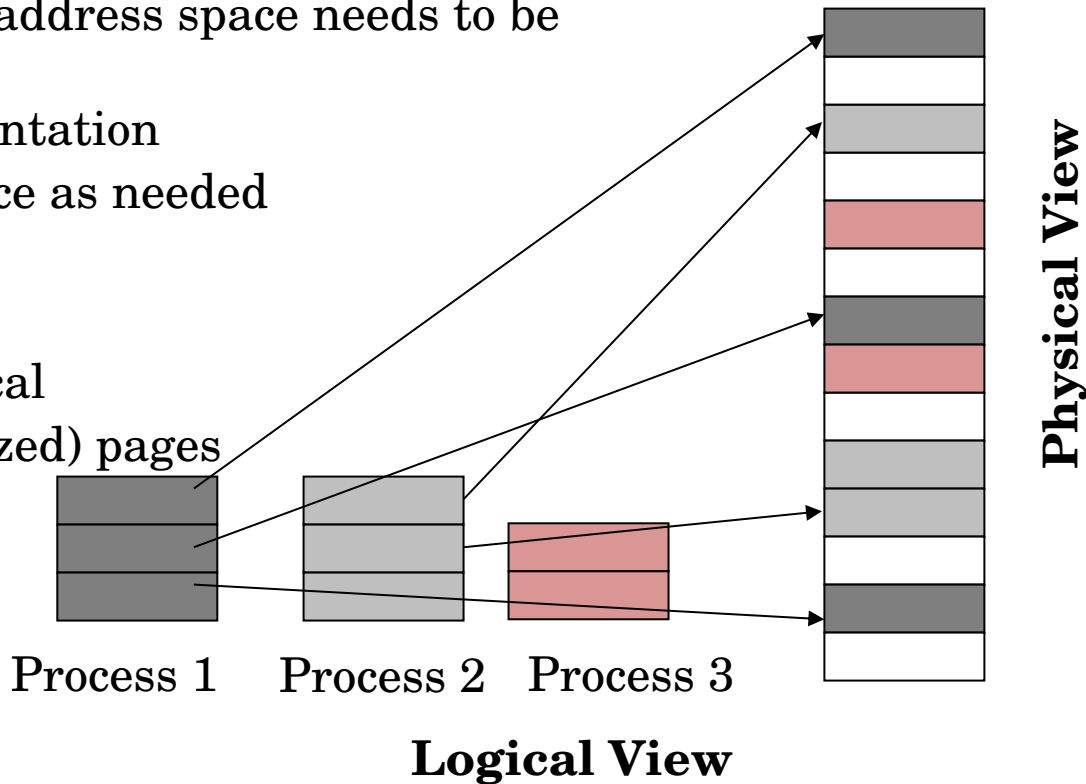
Goal: Remove requirement that address space needs to be contiguous

- Eliminates external fragmentation
- Ability to grow address space as needed

Idea:

Divide address spaces and physical memory into fixed-sized (same sized) pages

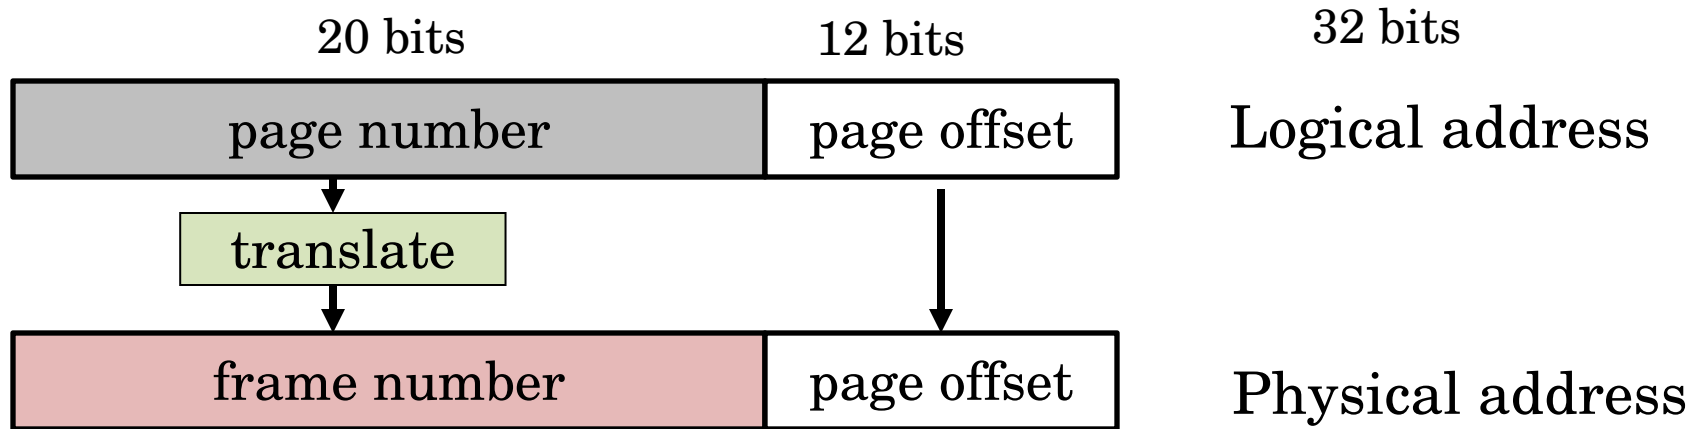
Size: 2^n , Example: 4KB



TRANSLATION OF PAGE ADDRESSES

How to translate logical address to physical address?

- High-order bits of address designate page number
- Low-order bits of address designate offset within page



No addition needed (unlike segmentation); just append bits correctly...