

MEMORY MANAGEMENT: SWAPPING TO DISK

Kai Mast

CS 537

Fall 2022

ANNOUNCEMENTS

- Project 2b will be released on Monday
 - Should be less work than P2a
 - Due on 10/24
- Midterm review session **in class** on 10/18


RECAP: SMALLER PAGETABLES

1. Segmented Pagetables
2. Multi-level Pagetables
 - Page the page tables
 - Page the page tables of page tables...
3. Inverted Pagetables

(Book Chapter 20)


MULTI-LEVEL PAGE TABLES

30-bit virtual address:



ADDRESS FORMAT FOR MULTI-LEVEL PAGING

30-bit address



How should logical address be structured? How many bits for each paging level?

Goal? Each page table fits within a page (assume page size = 4KB, 2^{12} bytes)

- number PTE * PTE size (assume 4bytes) = page size

How many page table entries can we fit on page?

- 4KB / 4bytes = 1K ($1024 = 2^{10}$) entries
- bits for selecting inner page = 10

Remaining bits for outer page: 30 bits - 12 bits - 10 bits = 8 bits

MULTI-LEVEL EXAMPLE

page directory

PPN	valid
0x3	1
-	0
-	0
-	0
-	0
-	0
-	0
-	0
-	0
-	0
-	0
-	0
-	0
-	0
-	0
-	0
-	0
0x92	1

page of PT (@PPN:0x3)

PPN	valid
0x10	1
0x23	1
-	0
-	0
0x80	1
0x59	1
-	0
-	0
-	0
-	0
-	0
-	0
-	0
-	0
-	0
-	0
-	0
-	0

page of PT (@PPN:0x92)

PPN	valid
-	0
-	0
-	0
-	0
-	0
-	0
-	0
-	0
-	0
-	0
-	0
-	0
-	0
-	0
-	0
-	0
0x55	1
0x45	1

translate 0x01ABC

0x23ABC

translate 0x04000

0x80000

translate 0xFEE00

0x55ED0

20-bit address

outer page(4 bits)


inner page(4 bits)

page offset (12 bits)

PROBLEMS WITH TWO LEVELS?

Problem: page directories (outer level) may not fit in a page


40-bit address:



$\text{PageSize} / \text{PTE Size} = 4\text{KB} / 4 \text{ bytes} \Rightarrow \text{max } 1\text{K entries per level}$

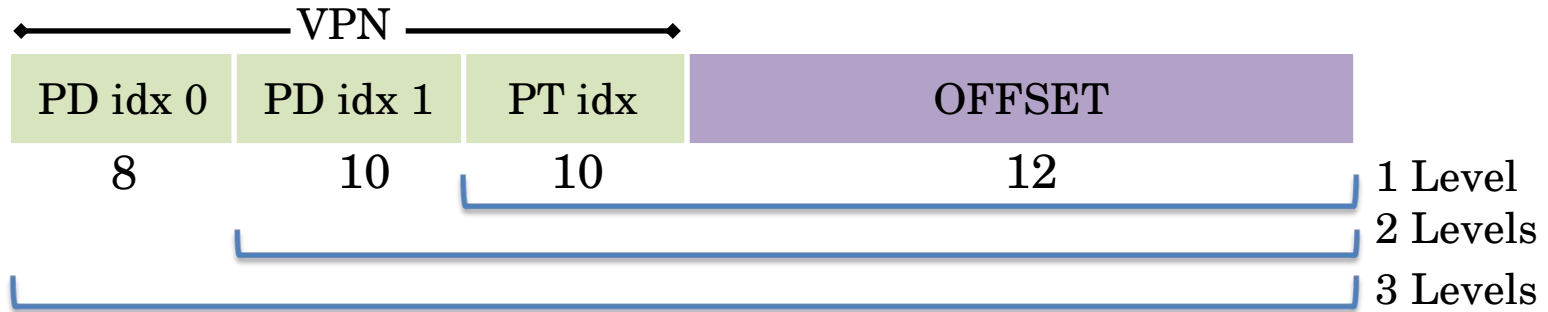
Solution:

- Split page directories into pieces
- Use another page directory to refer to the page directory pieces



MORE THAN TWO LEVELS OF PAGING

- Split page directories into pieces
- Use another page directory to refer to the page directory pieces



How large is virtual address space available with each level?

1 level:	<input type="text"/>
2 levels:	<input type="text"/>
3 levels:	<input type="text"/>

How many extra memory accesses?

FULL SYSTEM WITH TLBS

On TLB miss: **look-ups across multiple levels are more expensive**

Assume 3-level page table and 256-byte pages

Assume 16-bit addresses

Assume ASID of current process is 211

ASID	VPN	PFN	Valid
211	0xbb	0x91	1
211	0xff	0x23	1
122	0x05	0x91	1
211	0x05	0x12	0

How many physical memory accesses for each instruction? (ignore ops changing TLB)

(a) 0xAA10: movl 0x1111, %edi 0xaa: (TLB miss -> 3 for addr trans) + 1 instr fetch

Total: 8

0x11: (TLB miss -> 3 for addr trans) + 1 movl

(b) 0xBB13: addl \$0x3, %edi 0xbb: (TLB hit -> 0 for addr trans)

Total: 1

+ 1 instr fetch from 0x9113

(c) 0x0519: movl %edi, 0xFF10 0x05: (TLB miss -> 3 for addr trans) + 1 instr fetch

Total: 5

0xff: (TLB hit -> 0 for addr trans) + 1 movl into 0x2310

OTHER APPROACHES

~~1. Segmented Pagetables~~

~~2. Multi-level Pagetables~~

~~— Page the page tables~~

~~— Page the page tables of page tables...~~

3. Inverted Pagetables

INVERTED PAGE TABLES

Observation

- We only need entries for virtual pages w/ valid physical mappings (assuming there is no swapping; see later...)
- Have entries based on physical pages (**inverted**)

Naïve approach

- Search through data structure $\langle \text{PFN}, \text{VPN} + \text{ASID} \rangle$ to find match
- Too much time to search entire structure

Faster

- Find possible matches entries by **hashing** VPN+ASID
- Smaller number of entries to search for exact match

TLB still manages most cases

- Managing inverted page table requires software-controlled TLB

SUMMARY: SMALLER PAGE TABLES

Problem: Simple linear page tables require too much contiguous memory

Many options exist for efficiently organizing page tables

If OS traps on TLB miss, OS can use any data structure

- Inverted page tables (hashing)

If Hardware handles TLB miss, page tables must follow specific format

- Multi-level page tables used in x86 architecture
- Each page table fits within a page

Next Topic: What if allocated virtual memory does *not* fit into physical memory?


SWAPPING MEMORY

(Book Chapter 21)

PAGING USES: SHARED MEMORY

- Exploit level of indirection between VA and PA
- Regions of two separate processes' address spaces map to the same physical pages
 - read/write: access to share data
 - execute: shared libraries!
- Separate PTEs per process
 - Easy to provide different access privileges

SHARED MEMORY VISUALIZATION




PAGING USES: COPY-ON-WRITE


Copy-on-write (COW)

- Do not copy all pages; rather share the mappings
- Allows sharing of pages before changes are made
- Child process
 - Read-only mapping
 - On a write, protection fault occurs; copy page and resume
- Allows doing `fork()` very efficiently

COW – BEFORE MODIFICATION



COW – AFTER MODIFICATION

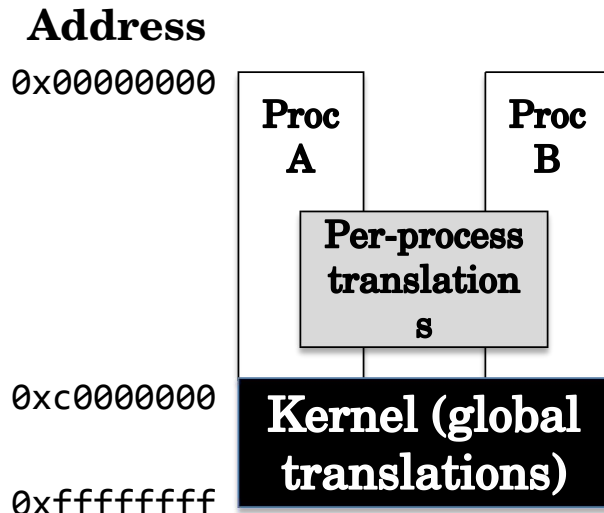


SIDE-NOTE: USER VS KERNEL ADDRESSES

- Low region of address space is private per-process memory
- High region reserved for kernel use and has same translations for all processes

Challenge: How to avoid accessing kernel memory

- Privileged bit in PTE marks high region as only accessible when in privileged mode



SWAPPING MOTIVATION

OS Goal: Allow execution even when not enough physical memory available

- Single process with very large address space
- Multiple processes whose total address spaces are large

Process execution should be independent of the amount of physical memory


- Correctness, if not performance

Virtual memory: OS provides illusion of more physical memory

Why does swapping work?

- Relies on key properties of user processes (workload) and machine architecture (hardware)

WORKLOAD HELP



What is in a program's code?

Code: many large libraries, some of which are rarely/never used

How to avoid wasting physical pages to store rarely used virtual pages?

Paging in and out

WORKLOAD: LOCALITY OF REFERENCE

Leverage **locality of reference** within processes


- **Spatial**: reference memory addresses **near** previously referenced addresses
- **Temporal**: reference memory addresses that have referenced in the **past**
- Processes spend majority of time in a small portion of their code
 - Estimate: 90% of time in 10% of code

Implication:

- Process only uses small amount of address space at any moment
- Only small amount of address space must be resident in physical memory for reasonable performance

HARDWARE HELP: MEMORY HIERARCHY

Leverage **memory hierarchy** of machine architecture
Each layer acts as “backing store” for layer above



What entity controls the placement of data into each layer?

SWAPPING INTUITION

Idea: OS keeps unreferenced pages on disk (HDD or SSD)

- Slower, cheaper backing store than memory

Process can run even when not all its pages are loaded into main memory

- OS and hardware cooperate to make large disk seem like memory
- Same behavior (correctness) as if all of address space in main memory

Requirements:

- OS must have a **mechanism** to identify location of each page in address space in memory or on disk
- OS must have a **policy** for determining which pages live in memory and which on disk

SWAPPING MECHANISMS

Each page in virtual address space maps to one of three locations:

- Physical main memory: Small capacity, fast performance, expensive cost
- Disk (backing store): Large capacity, slow performance, inexpensive
- Nothing (error): Infinite size, infinitely fast, no cost

Extend page tables with extra bits: **present** and **dirty**

- Permissions (r/w), valid, present, dirty in every PTE
- Page in memory: present bit set
- Page on disk: present bit cleared
 - PTE points to block on disk
 - Causes trap into OS when page is referenced (page fault)
- Page in memory, but in-memory contents don't match those on disk: dirty bit set