

CONDITION VARIABLES AND THE PRODUCER/CONSUMER PROBLEM

Kai Mast

CS 537

Fall 2022

RECAP: LOCKS AND CONDITION VARIABLES

- Locks enable **mutual exclusion** of threads
 - Only at most one thread can execute a critical section
- Condition variables allow **ordering** of thread execution
 - Threads can wait() on the condition variable and will be woken up if

RECAP: CONDITION VARIABLES

Interface

- `wait()`: block current thread until woken up by another thread
- `signal/broadcast()`: wake up (at least) one thread or wake up all threads

Common Errors

- No (re-)checking of state before calling `wait`
- Not holding the lock while changing shared state
- `if` instead of `while`; does not protect against spurious wake-ups

Spurious Wake-ups

- Problem: Signal might wake up more than one waiter
- Or waiters get woken up without even when there was no signal (less common)
- Why? Might have been more efficient to implement

RECAP: JOIN() IMPLEMENTATION

```
void thread_exit(thread_t *t) {  
    mutex_lock(&t->mutex);  
    t->done = 1;  
    cond_signal(&t->cond);  
    mutex_unlock(&t->mutex);  
}
```

```
void thread_join(thread_t *t) {  
    mutex_lock(&t->mutex);  
    while (t->done == 0) {  
        cond_wait(&t->cond, &t->mutex);  
    }  
    mutex_unlock(&t->mutex);  
}
```

Why do we need the done-variable?

- Thread might already have terminated before we call thread_join()

Why do we need the while-loop?

- To prevent against spurious wake-ups
- Generally, always (re-)check state after acquiring a lock

A MORE COMPLEX EXAMPLE: UNIX PIPES

What happens when we do this?

```
prompt> cat myfile.txt | grep "cs537"
```

- OS maps `cat`'s standard output to a pipe
- OS maps `grep`'s standard input to the same pipe
- Why is this useful?
 - Allows performing **complex tasks** using a combination of **simple tools**
- Pipe is a limited size buffer provided by the OS
 - Why limited? Dangerous to let process fill an unlimited kernel buffer

A MORE COMPLEX EXAMPLE: UNIX PIPES

```
prompt> cat myfile.txt | grep "cs567"
```

- A pipe can have multiple readers and writers
- Internally, there is a finite-size, circular buffer
 - We'll see later what "circular" means
- Writers add data to the buffer: May have to wait if buffer is full
- Readers remove data from the buffer: May have to wait if buffer is empty

MORE GENERAL: PRODUCERS & CONSUMERS

Producers generate data (e.g., pipe writers)

Consumers grab data and process it (e.g., pipe readers)

Producer/consumer problems are frequent in systems (e.g., web servers)

General strategy use condition variables to:

- Make producers wait when buffers are full
- Make consumers wait when buffers are empty

Handle case where producer and consumer work at much **different speeds**

SINGLE PRODUCER/CONSUMER

We start with an easy case:

- One producer thread
- One consumer thread
- At most one item that can be produced or consumed

Keep track of array state using shared state

- `num_full` indicates the available items to “consume”
- Can only be 0 or 1

SINGLE PRODUCER/CONSUMER

```
void* producer(void *arg) {  
    while (!done) {  
        mutex_lock(&m);  
        while (num_full > 0)  
            cond_wait(&cond, &m);  
        fill_buffer();  
        num_full += 1;  
        cond_signal(&cond);  
        mutex_unlock(&m);  
    }  
}
```

```
void* consumer(void *arg) {  
    while (1) {  
        mutex_lock(&m);  
        while (num_full == 0)  
            cond_wait(&cond, &m);  
        use_buffer();  
        num_full -= 1;  
        cond_signal(&cond);  
        mutex_unlock(&m);  
    }  
}
```

Will the producer be stuck waiting for mutex lock()?

What happens to consumer when producer calls signal?

PRODUCER/CONSUMER WITH MULTIPLE BUFFERS

Now a slightly more complicated case:

- One producer thread
- One consumer thread
- An array of multiple buffers that can be filled or consumed

Keep track of array state using shared state

- `num_full` indicates the available items to “consume”
- `max` indicates the size of the array
- `num_full` can be 0 or any integer $\leq \text{max}$

PRODUCER/CONSUMER WITH MULTIPLE BUFFERS



```
void fill_next_buffer(int value) {  
    buffer[fill_ptr] = value;  
    fill_ptr = (fill_ptr + 1) % max;  
}
```

```
int use_next_buffer() {  
    int tmp = buffer[use_ptr];  
    use_ptr = (use_ptr + 1) % max;  
    return tmp;  
}
```

- `fill_ptr` tracks the current position of the producer
- `use_ptr` tracks the current position of the consumer
- When reaching the end of the buffer, pointers wrap around

PRODUCER/CONSUMER WITH MULTIPLE BUFFERS

```
void* producer(void *arg) {  
    while (!done) {  
        mutex_lock(&m);  
        while (num_full == max)  
            cond_wait(&cond, &m);  
        fill_next_buffer();  
        num_full += 1;  
        cond_signal(&cond);  
        mutex_unlock(&m);  
    }  
}
```

```
void* consumer(void *arg) {  
    for (1) {  
        mutex_lock(&m);  
        while (num_full == 0)  
            cond_wait(&cond, &m);  
        use_next_buffer();  
        num_full -= 1;  
        cond_signal(&cond);  
        mutex_unlock(&m);  
    }  
}
```

Does this behave correctly?

Yes

How much concurrency does this allow?

Only one thread can fill or use a buffer at a time

PRODUCER/CONSUMER WITH MULTIPLE BUFFERS

```
void* producer(void *arg) {  
    while (!done) {  
        mutex_lock(&m);  
        while (num_full == max)  
            cond_wait(&cond, &m);  
        mutex_unlock();  
        fill_next_buffer();  
        mutex_lock();  
        num_full += 1;  
        cond_signal(&cond);  
        mutex_unlock(&m);  
    }  
}
```

```
void* consumer(void *arg) {  
    for (1) {  
        mutex_lock(&m);  
        while (num_full == 0)  
            cond_wait(&cond, &m);  
        mutex_unlock();  
        use_next_buffer();  
        mutex_lock();  
        num_full -= 1;  
        cond_signal(&cond);  
        mutex_unlock(&m);  
    }  
}
```

Why can we release the lock here?

- num_full guards which buffers will be filled or used
- We only access num_full while holding the lock

MULTIPLE PRODUCERS/CONSUMERS

Another slightly more complicated case:

- **Multiple** producer threads
- **Multiple** consumer threads
- At most one item that can be produced or consumed

Keep track of array state using shared state

- `num_full` indicates the available items to “consume”
- Can only be 0 or 1

MULTIPLE PRODUCERS/CONSUMERS

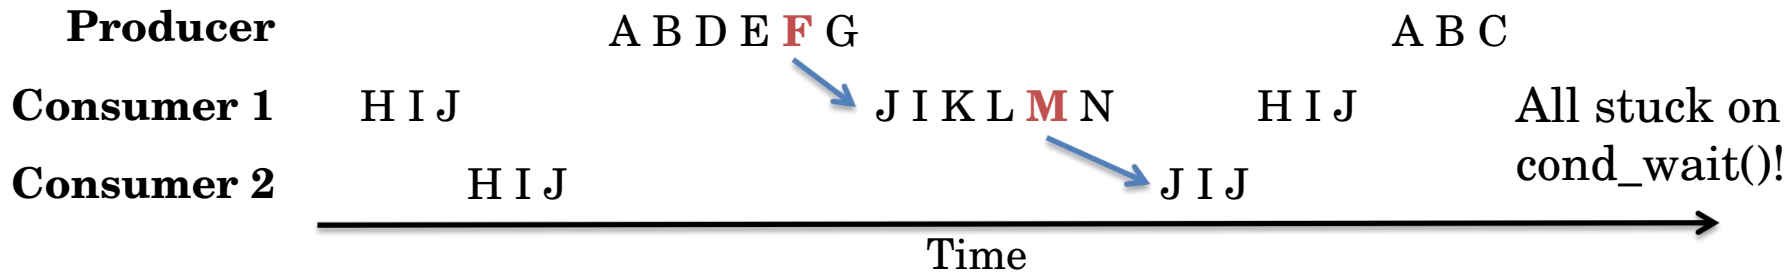
```
void* producer(void *arg) {
    while (1) {
        mutex_lock(&m);
        while (num_full > 0)
            cond_wait(&cond, &m);
        fill_buffer();
        num_full += 1;
        cond_signal(&cond);
        mutex_unlock(&m);
    }
}
```

A
B
C
D
E
F
G

```
void* consumer(void *arg) {
    while (1) {
        mutex_lock(&m);
        while (num_full == 0)
            cond_wait(&cond, &m);
        use_buffer();
        num_full -= 1;
        cond_signal(&cond);
        mutex_unlock(&m);
    }
}
```

H
I
J
K
L
M
N

What is the potential problem here?



SIMPLE SOLUTION: USE BROADCAST

```
void* producer(void *arg) {  
    while (!done) {  
        mutex_lock(&m);  
        while (num_full > 0)  
            cond_wait(&cond, &m);  
        fill_buffer();  
        num_full += 1;  
        cond_broadcast(&cond);  
        mutex_unlock(&m);  
    }  
}
```

```
void* consumer(void *arg) {  
    while (1) {  
        mutex_lock(&m);  
        while (num_full == 0)  
            cond_wait(&cond, &m);  
        use_buffer();  
        num_full -= 1;  
        cond_broadcast(&cond);  
        mutex_unlock(&m);  
    }  
}
```

Does this behave correctly?

Yes

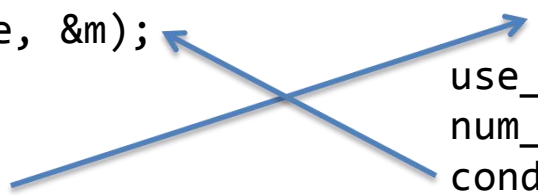
Is this an ideal solution?

No, scales poorly with number of threads

BETTER: USE TWO CONDITION VARIABLES

```
void* producer(void *arg) {  
    while (!done) {  
        mutex_lock(&m);  
        while (num_full > 0)  
            cond_wait(&consume, &m);  
        fill_buffer();  
        num_full += 1;  
        cond_signal(&produce);  
        mutex_unlock(&m);  
    }  
}
```

```
void* consumer(void *arg) {  
    while (1) {  
        mutex_lock(&m);  
        while (num_full == 0)  
            cond_wait(&produce, &m);  
        use_buffer();  
        num_full -= 1;  
        cond_signal(&consume);  
        mutex_unlock(&m);  
    }  
}
```



produce-condition is only signaled by producers and only wakes up consumers

consume-condition is only signaled by consumers and only wakes up producers

CV BEST PRACTICES REVISITED

- Do not use one condition variable across multiple mutexes
- But, you can use **multiple condition variables for one mutex**
 - Make sure each condition variable has a designated “role”
- As with locks, more condition variables make the code **more complex** but, potentially, **increase performance**
 - Find a good middle ground between the two

ANOTHER EXAMPLE: MEMORY MANAGEMENT

- In this example, heap memory is limited
 - Threads might block on `alloc()`
- Very simplified: we only keep track of an integer value; not the actual memory mapping
- Can be generalized to other such non-uniform producer/consumer patterns

MEMORY MANAGEMENT EXAMPLE

```
void* allocate(int size) {  
    mutex_lock(&m);  
    while (bytes_left < size)  
        cond_wait(&c, &m);  
    // get mem from heap  
    void *ptr = ...;  
    bytes_left -= size;  
    mutex_unlock(&m);  
    return ptr;  
}
```

```
void free(void *ptr, int size) {  
    mutex_lock(&m);  
    bytes_left += size;  
    cond_signal(&c);  
    mutex_unlock(&m);  
}
```

Problem?

- Thread that gets woken up might need more memory than became available

MEMORY MANAGEMENT FIXED

```
void* allocate(int size) {  
    mutex_lock(&m);  
    while (bytes_left < size)  
        cond_wait(&c, &m);  
    // get mem from heap  
    void *ptr = ...;  
    bytes_left -= size;  
    mutex_unlock(&m);  
    return ptr;  
}
```

```
void free(void *ptr, int size) {  
    mutex_lock(&m);  
    bytes_left += size;  
    cond_broadcast(&c);  
    mutex_unlock(&m);  
}
```

- All threads will be woken up
- If there is sufficient memory for a thread, that thread will be able to make progress
- A more efficient, but more complicated, solution probably exists