

PERSISTENCE: RAID

Kai Mast

CS 537

Fall 2022

HARD DISK ACCESS SPEED

To perform an I/O operation

- **seek:** move disk arm to correct track
- **wait (rotation):** wait for sector to rotate under arm
- **transfer:** read/write data

Example: Read one sector (512B)

- avg. seek: 7ms
- avg. rotate: 3ms
- avg. transfer: ~0ms (200Mb/s)
- throughput is $512\text{b}/10\text{ms} \approx 50\text{kb/s}$

Why so slow?

- Random I/O are dominated by seek and rotation
- Sequential accesses can be much faster

DISK SCHEDULING

Goal: Maximize throughput by minimizing seek and rotation times

Controller keeps track of multiple outstanding operations

- Duration of each access is known
- **Shortest-Seek First (SSF)**: To maximize throughput pick next operation with smallest seek time
- Can be implemented by the OS as **nearest block first**

Problems with SSF

- Does not account for rotation
- Disk arm might stay on same track for a long time
- Some operations could starve

ELEVATOR SCAN ALGORITHM

SCAN or “Elevator”: Moves from one end to the platter to the other and back

F-SCAN for “Freeze”: Executes a fixed number of operations in one batch

- Executions other operations later
- Fairer to requests that are on other parts of the platter

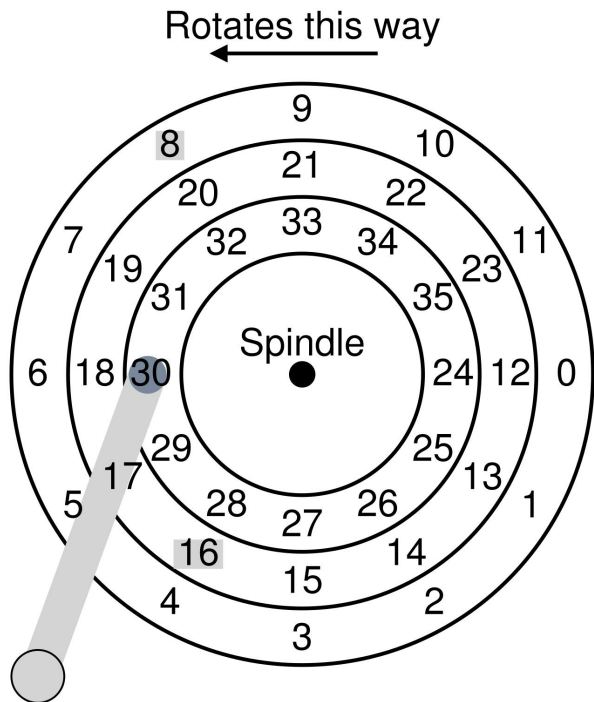
C-SCAN for “Circular”: Only moves into one direction

- Does not favor middle tracks

Problems:

- Does not take rotation time into account, only seek
- Better starvation free “Shortest-Job First” algorithms exist

SHORTEST-POSITIONING TIME FIRST



On modern devices seek and rotation times are roughly equivalent

- We should minimize both

Example

- We just accessed track 30
- SCAN would pick next
- Better to access next

Needs knowledge about the platter characteristics

- Best to implement in the disk



RAID

(Book Chapter 38)

WHY MULTIPLE DISKS?

Sometimes we want to use many disks — why?

- Capacity
- Reliability
- Performance


Challenge:

Most file systems manage only one disk
(assume linear array of blocks)

SOLUTION 1: JBOD

JBOD: **J**ust a **B**unch **O**f **D**isks

Application stores different files on different file systems



Disadvantages:


- Application must manage multiple devices
- Not portable across different system configurations

SOLUTION 2: RAID

RAID: **R**edundant **A**rray of **I**nexpensive **D**isks

RAID is

- transparent
- easy to deploy



Logical disk gives

- capacity
- performance
- reliability

Logical disk abstracts away underlying physical disks

WHY INEXPENSIVE DISKS?

Alternative to RAID


- Buy one high performance, highly reliable, high-capacity disk

RAID Approach

- Economies of scale! Commodity disks cost less per byte
- Can buy **many** commodity H/W components for same price as few high-end components
- Write software to build high-quality logical devices from many cheap devices

GENERAL STRATEGY: MAPPING

Build fast, large disk from smaller disks



RAID MAPPING

How should RAID map logical block addresses to physical block addresses?


- Some similarity to virtual memory

1) **Dynamic** mapping (page table approach)

- Logical x sometimes maps to physical y and sometimes z
- Use data structure (array, hash table, tree)

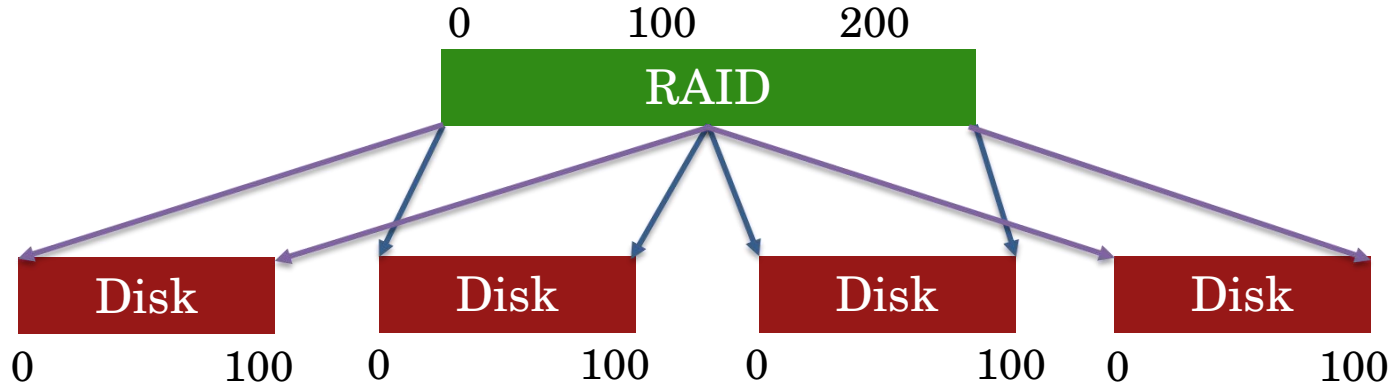
2) **Static** mapping (RAID approach)

- Logical x always maps to physical
- Uses simple math; avoids extra look-ups



GENERAL STRATEGY: REDUNDANCY

Add even more disks for **reliability**



REDUNDANCY


How many physical **copies** should RAID keep for every logical block?

Increase number of copies:

- Improves **reliability** (and maybe **performance**)

Decrease number of copies

- Improves **space efficiency**



REDUNDANCY: FAILURE MODEL

Assume disks are **fail-stop**

- Will either execute reads/writes correctly or do nothing
- System knows when disk fails

Most disk failures can be turned into fail-stop

- e.g., data corruption can be detected using error correction codes

REASONING ABOUT RAID

Pick one of each

1) RAID Level

System for mapping logical to physical blocks

2) Workload

Types of reads/writes issued by applications (sequential vs. random)

3) Metric

Capacity, reliability, performance

1) RAID DECISIONS

Which logical blocks map to which physical blocks on disks?

How to use redundancy (if any)?

Different **RAID levels** make different trade-offs

- RAID 0: Striping

- RAID 1: Mirroring

- RAID 4: Parity

- RAID 5: Rotated Parity

2) WORKLOADS

Reads

- One operation (for latency)
- Steady-state I/O (for throughput or bandwidth)
 - Sequential or Random

Writes

- One operation (for latency)
- Steady-state I/O (for throughput or bandwidth)
 - Sequential
 - Random

3) METRICS

Capacity: how much space is available to higher levels?

Reliability: how many disks can RAID safely lose? (assume fail stop!)

Performance: how long does each workload take?

Normalize each to characteristics of one disk

N := total number of physical disks

C := capacity of each disk (e.g., 500 GB)


S := sequential throughput of each disk (e.g., 100 MB/s)

R := random throughput of each disk (e.g., 5 MB/s)

D := latency of one small I/O operation

RAID-0: STRIPING

Optimizes for **capacity**. Does not provide **redundancy**.



RAID-0: 4 DISKS

	Disk 0	Disk 1	Disk 2	Disk 3
	0	1	2	3
stripe	4	5	6	7
	8	9	10	11
	12	13	14	15

Given logical address A, find:

Disk = $A \% \text{disk_count}$

Offset = $A / \text{disk_count}$

REAL SYSTEMS: CHUNK SIZE

	Disk 0	Disk 1	Disk 2	Disk 3
Chunk size = 1	0	1	2	3
	4	5	6	7
	8	9	10	11
	12	13	14	15

Simplification: assume chunk size of 1

	Disk 0	Disk 1	Disk 2	Disk 3
Chunk size = 2	0 1	2 3	4 5	6 7
	8 9	10 11	12 13	14 15

stripe: 8 10 12 14

RAID-0: ANALYSIS

What is the capacity? $N * C$

How many disks can fail (no loss)? 0

Latency? D

Sequential Throughput? $N*S$

Random Throughput? $N*R$

Adding more disks improves throughput, but not latency!

N := number of disks

C := capacity of 1 disk


S := sequential throughput of 1 disk

R := random throughput of 1 disk

D := latency of one small I/O operation


Disk 0	Disk 1	Disk 2	Disk 3
0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

RAID-1: MIRRORING



Keep two copies of every block

RAID-1: MIRRORING



How many disks can fail without losing any data?

RAID-1 can always handle 1 disk failure

RAID-1 WITH STRIPING

Disk 0	Disk 1	Disk 2	Disk 3
0	0	1	1
2	2	3	3
4	4	5	5
6	6	7	7

Combines mirroring and striping

- Capacity is $N \cdot C / 2$
- Can handle at least 1 and up to $N/2$ failures
 - e.g., 1 and 3 can fail, but not 2 and 3
- Also called RAID 10 or RAID 1+0

RAID-1: ANALYSIS

What is the capacity?	$N/2 * C$
How many disks can fail?	1 (or maybe $N / 2$)
Read Latency?	D
Write Latency?	D

N := number of disks

C := capacity of 1 disk

S := sequential throughput of 1 disk

R := random throughput of 1 disk

D := latency of one small I/O operation

Disk 0	Disk 1	Disk 2	Disk 3
0	0	1	1
2	2	3	3
4	4	5	5
6	6	7	7

RAID-1: THROUGHPUT

What is the steady-state throughput for

- random reads? $N * R$
- random writes? $N/2 * R$
- sequential writes? $N/2 * S$
- sequential reads? $N/2 * S$

If a fully sequential read is split across all four disks, each disk would spend half its time spinning to the next location

Disk 0	Disk 1	Disk 2	Disk 3
0	0	1	1
2	2	3	3
4	4	5	5
6	6	7	7

SIDE ISSUE: SYSTEM CRASHES

RAID 1

Block	Disk 0	Disk 1
0	A	A
1	X	B
2	C	C
3	D	D

System crashes can happen due to bugs or power loss

- Requires reboot!

Application writes “X” to block 1

- Disk 0 writes to block 2 successfully
- System crashes before Disk 1 is done writing to block 2

Problem: After reboot, how to tell which data is right?

CRASHES: H/W SOLUTION

Consistent-Update Problem:

We want writes on both/all disk to be **atomic**

Solution: Use non-volatile RAM in RAID controller

- Can replay to ensure all copies are updated
- Software RAID controllers (e.g., Linux md) don't have this option

RAID-4 STRATEGY

RAID-4: Compromise between RAID-0 and RAID-1

Use **one** disk for **parity** (form of redundancy, but not full replication)

In algebra: Equation with N variables and $N-1$ are known, can often solve for unknown

Treat sectors across disks in a stripe as equation

Data on bad disk is the unknown in equation