

MEMORY MANAGEMENT

Kai Mast

CS 537

Fall 2022

ANNOUNCEMENTS

- P1B is out
- My office hours are only on Tuesdays (not Thu) from now on
 - You can always schedule one for another time!
- I will try to upload slides the day before the lecture
 - But there might be minor changes and fixes later

RECAP: SCHEDULING

Scheduler: Decides which job to run

Job: A CPU burst of a process

Turnaround Time: Time from arrival to completion

Response Time: Time from arrival to first time running

Starvation: Job never gets scheduled

SCHEDULING POLICIES

FIFO:

- The job that arrives first, runs first


Shortest-Job First:

- Run the shortest ready job
- Best for minimizing average turnaround time

Round-Robin:

- At the end of every time slice, preempt the current process and switch to the next one
- Improves response time

MULTI-LEVEL FEEDBACK QUEUE (MLFQ)




Q7

Rule 1: If $\text{priority}(A) > \text{Priority}(B)$, A runs

Q6

Rule 2: If $\text{priority}(A) == \text{Priority}(B)$,
A & B run in Round-Robin

Q5




Rule 3: Processes start at top priority

Q3

Rule 4: If job uses whole slice, demote process
(longer time slices at lower priorities)

Q2



MLFQ INTUITION

Detect **interactive processes**

- Will frequently wait for user (and other) I/O


Schedule **batch processes** less frequently

- Usually not time sensitive

Reset levels regularly by **boosting** all processes to the highest level



- Allows adapting to workload changes

MLFQ EXAMPLE (NO I/O)




Short job arrives at $T=100$
and finishes at $T=120$

Legend

-  Long-Running Job
-  Short-Lived Job

MLFQ EXAMPLE (WITH I/O)



Interactive job does not use the full slice and stays at highest level

Legend



Batch Job (no I/O)



Interactive Job (has I/O)

INTERLUDE: PROCESS MANAGEMENT

PROCESS CREATION

Two ways to create a process

Option 1: Build a new process from scratch

Option 2: Copy an existing process and change it appropriately

OPTION 1: NEW PROCESS

Create new process with specified executable and state

- Load specified code and data into memory; Create empty call stack
- Create and initialize PCB (make look like context-switch)
- Put process on ready list

Advantages: No wasted work

Disadvantage: Difficult to setup process and to express all possible options

- Process permissions, where to write I/O, environment variables
- Example: Windows NT has call with 10 arguments

OPTION 2: COPY AND CHANGE

Copy existing process (fork) and change as needed (exec)

- `fork()`
 - Calling process (parent) creates a child process
 - Make copy of code, data, stack, and PCB of parent
 - Add new PCB to ready list
- `exec(const char *file)`
 - Replace current data and code segments with those in specified executable file

Advantages: Flexible, clean, simple

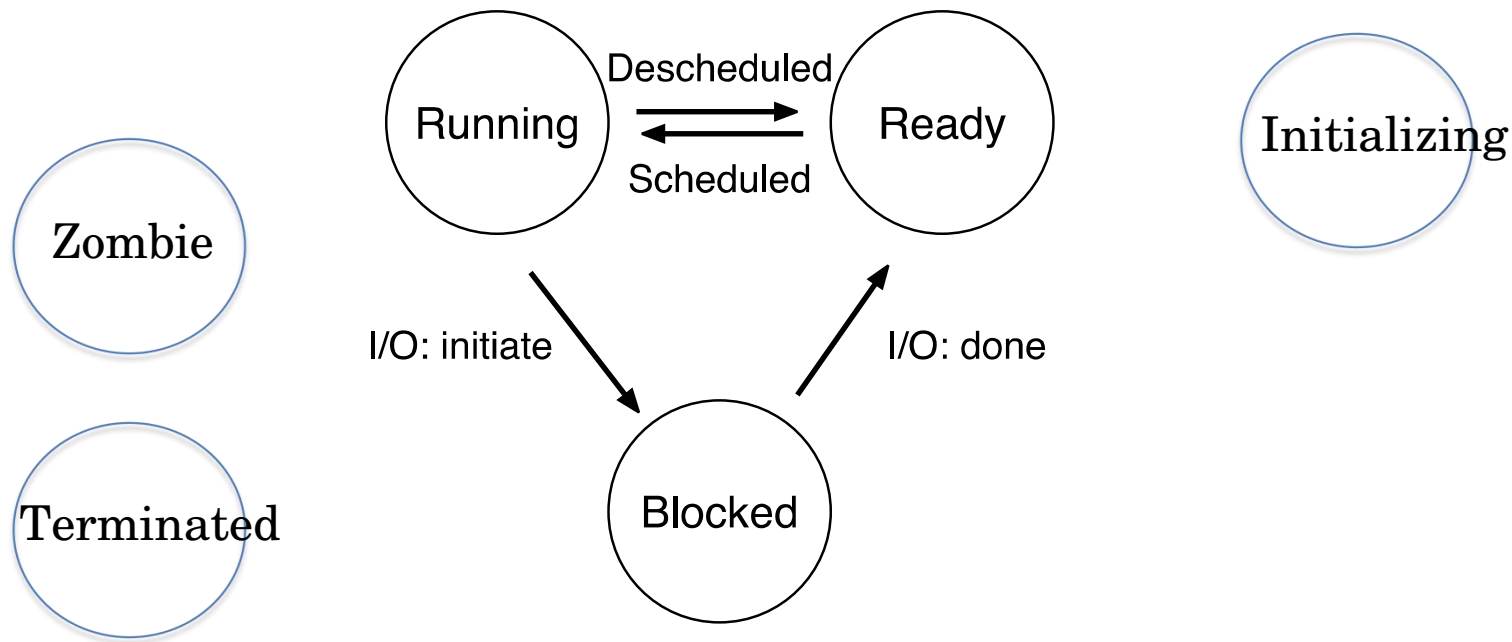
Disadvantages: Wasteful to perform copy and then overwrite of memory

UNIX SHELLS

```
while (true) {
    char *cmd = getcmd();
    int retval = fork();

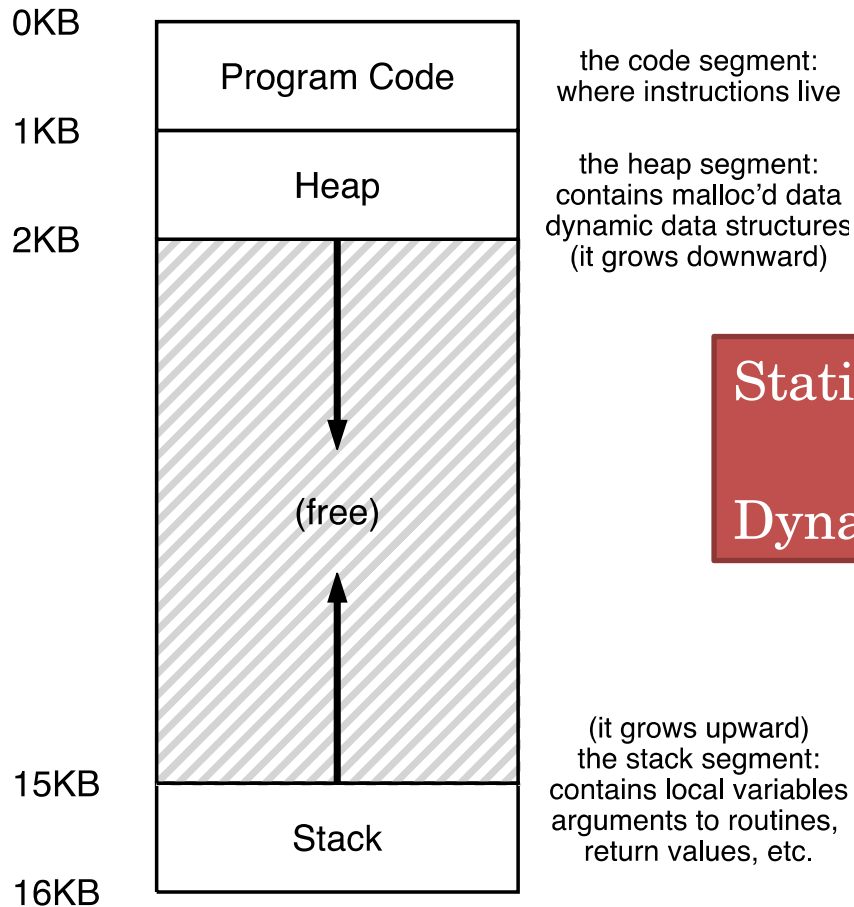
    if (retval == 0) {
        // This is the child process
        // Setup the child's process environment here
        // E.g., where is standard I/O, how to handle signals?
        exec(cmd);
    } else if (retval > 0) {
        // This is the parent process; Wait for child to finish
        int pid = retval;
        wait(pid);
    } else {
        // Handle errors here
    }
}
```

STATE TRANSITIONS EXTENDED



MEMORY MANAGEMENT


BACKGROUND: WHAT IS AN ADDRESS SPACE?



Static: Code and some global variables

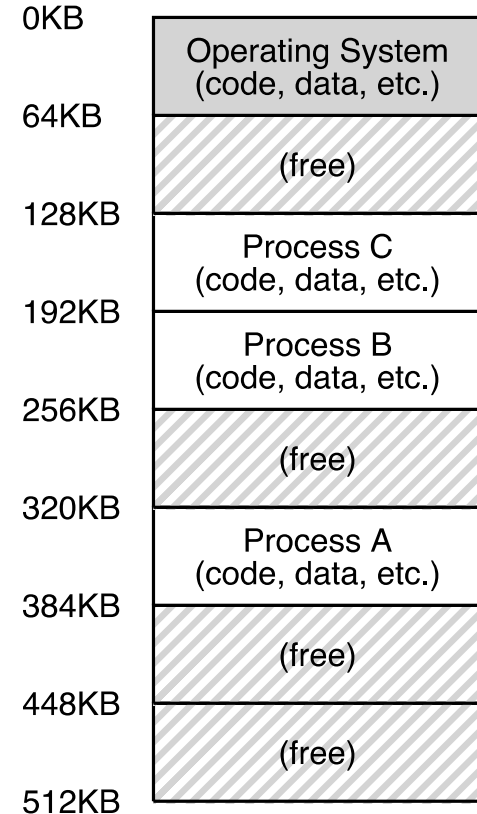
Dynamic: Stack and Heap

ABSTRACTION: ADDRESS SPACE



Each process has own set of addresses

How can the OS provide illusion of private (virtual) address space to each process?



REVIEW: MEMORY ACCESS (LOGICAL)

Initial %rip = 0x10 %rbp = 0x200

➡ 0x10: movl 0x8(%rbp), %edi
0x13: addl \$0x3, %edi
0x19: movl %edi, 0x8(%rbp)

%rbp is the base pointer:
points to base of current stack frame

%rip is instruction pointer (or program counter)

Fetch instruction at addr 0x10
Exec:

load from addr 0x208

Fetch instruction at addr 0x13
Exec:


addition (no memory access)

Fetch instruction at addr 0x19
Exec:

store to addr 0x208

5 total memory accesses

MOTIVATION FOR VIRTUALIZING MEMORY



First systems did not virtualize memory


Uniprogramming: One process runs at a time

Disadvantages

- Only one process be ready at a time
- Process can modify OS

Solutions?

TIME-SHARE MEMORY



- create P1
- swap out P1
- create P2
- swap out P2
- swap in P1
- ...

PROBLEMS WITH TIME SHARING

Ridiculously poor performance (disk are slow!)

Better Alternative:

- Space sharing of physical memory
- At same point in time, space of memory is divided across processes
- Remainder of solutions all use space sharing

STATIC RELOCATION

Rewrite each code segment before loading it in memory

- Pick static physical location for each process when started
- Each rewrite for different process uses different addresses and pointers
- Change jumps, loads of static data

STATIC RELOCATION

```
0x10:movl 0x8(%rbp), %edi
0x13:addl $0x3, %edi
0x19:movl %edi, 0x8(%rbp)
```

**Rewrite
for P1**

**Rewrite
for P2**

```
0x1010: movl 0x8(%rbp), %edi
0x1013: addl $0x3, %edi
0x1019: movl %edi, 0x8(%rbp)
```

```
0x3010: movl 0x8(%rbp), %edi
0x3013: addl $0x3, %edi
0x3019: movl %edi, 0x8(%rbp)
```

Process 1

4 KB

(free)

Program Code

Heap

(free)

8 KB

Stack

(free)

12 KB

Program Code

Heap

(free)

Process 2

16 KB

Stack

(free)

Why did the OS not rewrite the stack address?

STATIC RELOCATION: DISADVANTAGES

No Protection

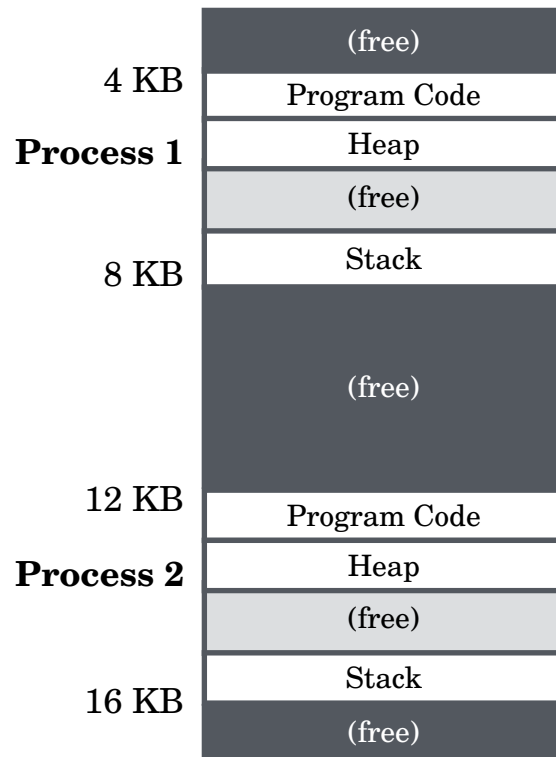
- Any process can modify memory of the OS or other processes

No Security

- All memory is visible to all processes

No Dynamic Allocation

- Cannot move address space after it has been placed
- May not have free space to allocate new process



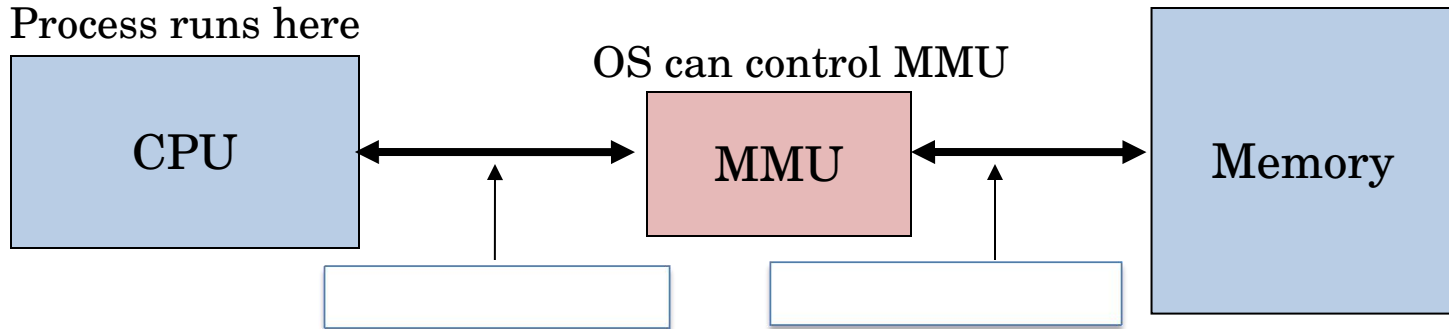
DYNAMIC RELOCATION

Goal: Protect processes from one another (and the OS from processes)

Requires hardware support: Memory Management Unit (MMU)

MMU dynamically changes process address at every memory reference

- Process generates **logical** or **virtual** addresses (in their address space)
- Memory hardware uses **physical** or **real** addresses



HARDWARE SUPPORT FOR DYNAMIC RELOCATION

Leverage privilege levels (introduced with system calls)

Non-privileged (user) mode: Processes run

- Perform **translation** of logical address to physical address

Privileged (protected, kernel) mode: OS runs

- When enter OS (trap, system calls, interrupts, exceptions)
- Allows privileged instructions to be executed
 - Can **manipulate contents of MMU**
- Allows **OS** to **access all of physical memory**

DYNAMIC RELOCATION USING BASE REGISTERS

Base Registers

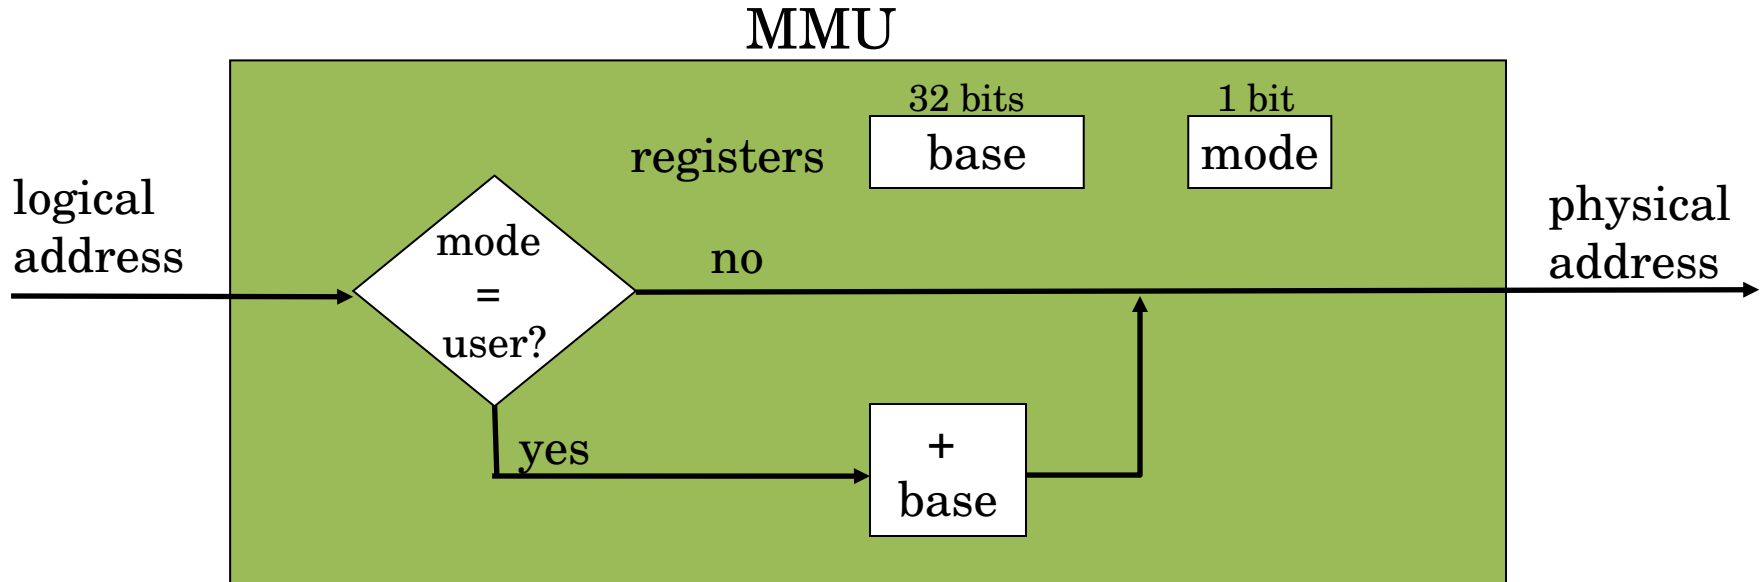
- Stored in the MMU
- Each process has different value in base register
- Set **by the OS** on **context switch**

Address Translation using Base Registers


- Add base register to virtual address on every memory access
- Dynamic relocation by changing value of base register!

BASE REGISTERS: IMPLEMENTATION

- Translation on every memory access of user process
- MMU adds base register to logical address to form physical address



PROTECTION WITH BASE REGISTERS?



Process	Base
P1	1024
P2	4096

Possible Execution

Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5096, R1
P1: load 1000, R1	load 2024, R1
P1: store 3072, R1	store 4096, R1

P1 can modify P2's address space!