

CONCURRENCY: LOCKS

Kai Mast

CS 537

Fall 2022

RECAP: THREADS

Threads enable **concurrency** within the same process

Each thread has its own

- Thread ID (TID)
- Set of registers, including **program counter** and **stack pointer**
- Stack for local variables and return addresses
 - There are now **multiple stacks in same address space**

EXAMPLE: CONCURRENT INCREMENT

Data

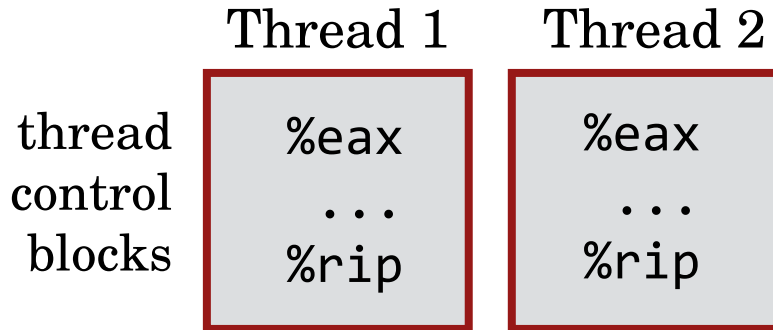
balance at 0x9cd4 = 100

Code (C)

```
balance = balance + 1;
```

Code (Assembly)

```
0x195  mov 0x9cd4, %eax
0x19a  add $0x1, %eax
0x19d  mov %eax, 0x9cd4
```



Registers are virtualized by OS;
Each thread “thinks” it has own

- Both threads run the above code
- Many possible **execution timelines** due to concurrency

PROBLEM: NON-DETERMINISM

Concurrency can lead to **non-deterministic results**

- Different results even with same inputs
- Race conditions

Whether bug manifests depends on CPU schedule!

How write concurrent programs: pretend **scheduler is malicious**

- Processes/threads can be stopped and resumed at any point
- Processes/threads might execute at different “speeds”

ATOMIC EXECUTION

We want instructions to execute as an uninterruptible group
That is, we want them to be **atomic**

```
mov 0x123, %eax  
add %0x1, %eax  
mov %eax, 0x123
```

More general: Need mutual exclusion for critical sections
if one thread is in critical section C, other threads cannot
(okay if other threads do unrelated work)

LOCKS

Goal: Provide mutual exclusion (**mutex**)

Allocate and Initialize

```
pthread_mutex_t mylock;  
pthread_mutex_init(&mylock, NULL);
```

Acquire (or Lock)

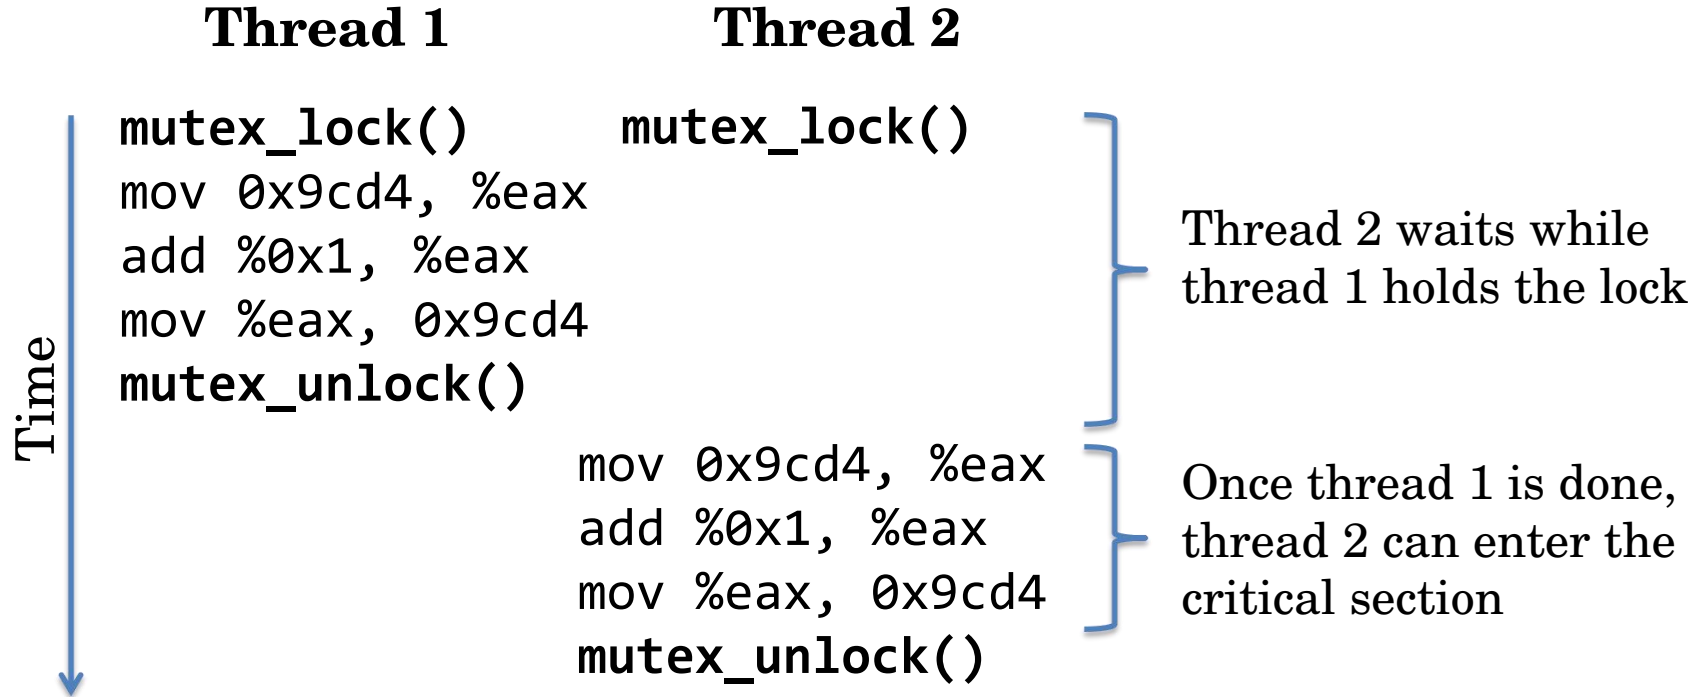
- Acquire exclusive access to lock
- Wait if lock is not available (some other process in critical section)
- Spin or block (relinquish CPU) while waiting
- **pthread_mutex_lock**(&mylock);

Release (or Unlock)

- Release exclusive access to lock; let another process enter critical section
- **pthread_mutex_unlock**(&mylock);

(Book Chapter 28)

EXECUTION TIMELINE WITH LOCKING



OTHER EXAMPLES

Consider multi-threaded applications that do more than incrementing a single integer value

Multi-threaded application with shared linked-list

- All concurrent:
 - Thread A inserting element a
 - Thread B inserting element b
 - Thread C looking up element c

EXAMPLE: SHARED LINKED LIST

```
typedef struct __node_t {  
    int key;  
    struct __node_t *next;  
} node_t;
```

```
typedef struct __list_t {  
    node_t *head;  
} list_t;
```

```
void list_init(list_t *L) {  
    L->head = NULL;  
}
```

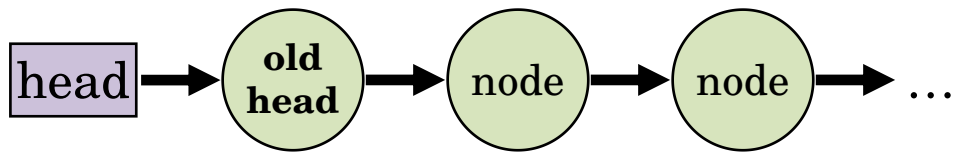
```
void list_insert(list_t *l, int key) {  
    node_t *new = malloc(sizeof(node_t));  
    assert(new);  
    new->key = key;  
    new->next = l->head;  
    l->head = new;  
}
```

```
int list_lookup(list_t *l, int key) {  
    for(node_t *curr = l->head; curr;  
        curr = curr->next) {  
        if (curr->key == key) {  
            return 1;  
        }  
    }  
    return 0;  
}
```

What can go wrong?

Find schedule that leads to problem?

LINKED-LIST RACE



```
void list_insert(list_t *l,  
                int key) {  
    node_t *new =  
        malloc(sizeof(node_t));  
    assert(new);  
    new->key = key;  
    new->next = l->head;  
    l->head = new;  
}
```

Thread 1

`new->key = key`

`new->next = L->head`

`l->head = new`

Thread 2

`new->key = key`

`new->next = L->head`

`L->head = new`

LOCKING LINKED LISTS

```
typedef struct __node_t {  
    int key;  
    struct __node_t *next;  
} node_t;
```

```
typedef struct __list_t {  
    node_t *head;  
} list_t;
```

```
void list_init(list_t *l) {  
    l->head = NULL;  
}
```

How to add locks?

```
void list_insert(list_t *l, int key) {  
    node_t *new = malloc(sizeof(node_t));  
    assert(new);  
    new->key = key;  
    new->next = l->head;  
    l->head = new;  
}
```

```
int list_lookup(list_t *l, int key) {  
    for(node_t *curr = l->head; curr;  
        curr = curr->next) {  
        if (curr->key == key) {  
            return 1;  
        }  
    }  
  
    return 0;  
}
```

A SIMPLE THREAD-SAFE LINKED LIST

```
typedef struct __node_t {
    int key;
    struct __node_t *next;
} node_t;

typedef struct __list_t {
    node_t *head;
    pthread_mutex_t lock;
} list_t;

void list_init(list_t *l) {
    l->head = NULL;
    pthread_mutex_init(&l->lock,
        NULL);
}
```

A SIMPLE THREAD-SAFE LINKED LIST (CONT.)

```
void list_insert(list_t *l, int key) {
    pthread_mutex_lock(l);
    node_t *new =
        malloc(sizeof(node_t));
    new->key = key;
    new->next = l->head;
    l->head = new;
    pthread_mutex_unlock(l);
}

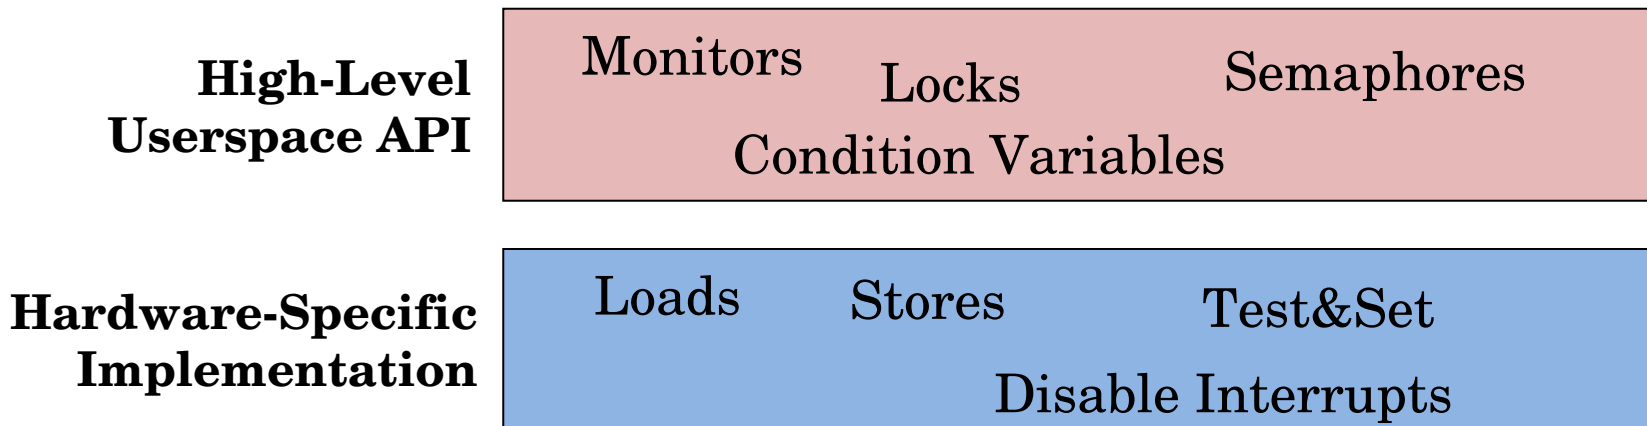
int list_lookup(list_t *l, int key) {
    pthread_mutex_lock(l);
    for (node_t *curr = l->head;
         curr!=NULL; curr = curr->next) {
        if (curr->key == key) {
            pthread_mutex_unlock(l);
            return 1;
        }
    }
    pthread_mutex_unlock(l);
    return 0;
}
```

What is the impact of critical section size?

SYNCHRONIZATION PRIMITIVES

- OS provides higher-level synchronization primitives
- Rely on hardware features if possible

Motivation: Build them once and get them right



LOCK IMPLEMENTATION GOALS

Correctness

- *Mutual exclusion*: Only one thread in critical section at a time
- *Progress*: If several simultaneous requests, must allow one to proceed

Performance:

- CPU is not used unnecessarily (e.g., no spin-waiting or busy-waiting)
- Fast to acquire lock if no contention with other threads (common case!)

Fairness: Each thread eventually gets a turn

- Wait times are *bounded* (starvation-free)
- Must eventually allow each waiting thread to enter (assuming others eventually exit)

APPROACH 1: DISABLE INTERRUPTS

```
void acquire_lock() {  
    disable_interrupts();  
}
```

```
void release_lock() {  
    enable_interrupts();  
}
```

Turn off interrupts while in a critical section

- Prevent scheduler from running another thread
- Code between interrupts executes atomically

Problems:

- Only works on systems with a single CPU
- OS yields control to the process
 - What if the process is malicious? What if it runs forever?
- Not fair: process might disable interrupts for a long time

APPROACH 2: USE LOAD AND STORE

```
void acquire_lock(lock_t *l) {  
    while (l->flag != 0) {}  
    l->flag = 1;  
}
```

```
void release_lock(lock_t *l) {  
    l->flag = 0;  
}
```

Approach

- Each lock has a single flag indicating its state
- acquire_lock **spins** until the lock is available

Does this work? What situation can cause this to not work?

PROBLEMS WITH APPROACH 2

Thread 1

l->flag = 1;

[critical section]

l->flag = 0;

Thread 2

while (l->flag != 0) {}

l->flag = 1;

[critical section]

l->flag = 0;

Thread 3

while (l->flag != 0) {}

l->flag = 1;

[critical section]

l->flag = 0;

When T1 releases lock, **both** T2 and T3 acquire it

THE NEED FOR HARDWARE SUPPORT

Software-based solutions exist but do not well on modern hardware

- See “Dekker’s and Peterson’s Algorithms” in the book

Better: Use a **hardware primitive** that can load and store atomically

- `test_and_set(&val, 1)` sets `val` to 1 and returns previous value
- Also called **atomic exchange**
- Specific instruction differs depending on CPU-architecture
 - e.g., `xchg` on x86

APPROACH 3: USE TEST-AND-SET

```
typedef struct __lock_t {  
    int flag;  
} lock_t;
```

```
void init(lock_t *lock) {  
    lock->flag = 0;  
}
```

```
void acquire(lock_t *lock) {  
    while (test_and_set(&lock->flag, 1) == 1) {}  
}
```

```
void release(lock_t *lock) {  
    lock->flag = 0;  
}
```

```
int test_and_set(int *var,  
                int new_value) {  
    int old_value = *var;  
    *var = new_value;  
    return old_value;  
}
```

OTHER POSSIBLE HARDWARE INSTRUCTIONS

```
int compare_and_swap(int *val, int cmp, int new)
```

- Set val to new if val==cmp
- Otherwise, does nothing
- Always returns val's previous value

```
int fetch_and_add(int *val)
```

- Increments val and returns its previous value

USE SPINNING WITH CAUTION

Consider the following scenario:

- Thread 1 is holding a lock for five time slices
- Ten other threads want to acquire the lock and spin
- The scheduler is Round-Robin with time slice length 10ms

What is the worst-case time spent on spinning? 500ms!

We need to avoid spinning as much as possible!

What could we do instead of spinning?

APPROACH 4: YIELD

Idea: Introduce a new system call `yield()` that voluntarily gives up CPU control

```
void acquire_lock(lock_t *l) {  
    while ((test_and_set(&l->flag, 1) == 1) {  
        yield();  
    }  
}  
  
void release_lock(lock_t *l) {  
    l->flag = 0;  
}
```

LOCK IMPLEMENTATION GOALS REVISITED

Yield-based locks (Approach 4)

- Correctness: Yes (only one thread can be in the CS)
- Performance: Better (still too many context switches)
- Fairness: No (scheduler does not know/respect wait order)

APPROACH 5: USE A QUEUE

Intuition: Two-step approach

- First lock a wait-queue and add current thread to it
 - Queue-lock is only held for a very short time
- When calling `release()`, pick the first thread from queue
- Needs to more system class
 - `park()` pauses/deschedules the current thread
 - `unpark(int thread_id)` resumes/schedules the specified thread

APPROACH 5: ACQUIRING LOCKS

```
typedef struct __lock_t {  
    int flag;  
    int guard;  
    queue_t *q;  
} lock_t;  
  
void init(lock_t *l) {  
    l->flag = 0;  
    l->guard = 0;  
    queue_init(l->q);  
}  
  
void acquire(lock_t *l) {  
    while(test_and_set(&l->guard,1) == 1)  
        ; //acquire guard lock by spinning  
    if (l->flag == 0) {  
        l->flag = 1; // lock is acquired  
        l->guard = 0;  
    } else {  
        queue_add(l->q, gettid());  
        l->guard = 0;  
        park();  
    }  
}
```

APPROACH 5: RELEASING LOCKS

```
typedef struct __lock_t {  
    int flag;  
    int guard;  
    queue_t *q;  
} lock_t;  
  
void init(lock_t *l) {  
    l->flag = 0;  
    l->guard = 0;  
    queue_init(l->q);  
}  
  
void release(lock_t *l) {  
    while(test_and_set(&l->guard,1) == 1)  
        ; //acquire guard lock by spinning  
    if (queue_empty(l->q)) {  
        // let go of lock; no one wants it  
        l->flag = 0;  
    } else {  
        // hold lock (for next thread!)  
        unpark(queue_remove(m->q));  
        l->guard = 0;  
    }  
}
```

SUMMARY

- Locks are a fundamental building block for concurrent applications and data structures
- Efficient locking mechanism require hardware and OS support
 - `compare_and_swap()`, atomic exchange, etc.
 - `yield()`, `park()`, etc.